

Package ‘ROOT’

November 19, 2025

Title Identifying Underrepresented Subpopulations With Interpretable Trees

Version 0.0.0.9000

Description ROOT (Rashomon set of Optimal Trees) is a framework for learning interpretable binary weight functions represented as sparse decision trees. It constructs a Rashomon set of near-optimal trees and extracts a characteristic tree to summarize patterns. Given trial and target data, the package identifies trial subpopulations that contribute disproportionately to the variance of the target treatment-effect estimate (underrepresented groups).

License MIT + file LICENSE

Encoding UTF-8

Roxxygen list(markdown = TRUE)

RoxxygenNote 7.3.3

Suggests mlbench,
testthat (>= 3.0.0),
knitr,
rmarkdown,
ragg

Config/testthat.edition 3

Imports MASS,
rpart,
gbm,
stats,
withr,
rpart.plot

VignetteBuilder knitr

Depends R (>= 3.5)

LazyData true

Contents

characterize_tree	2
characterizing_underrep	3
choose_feature	5
estimate	5
estimate_dml	6
estimate_dml_single	7

estimate_single	8
gen_S	9
gen_T	9
gen_XY	10
get_data	11
loss_from_objective	12
midpoint	12
objective_default	13
objective_if	13
plot.characterizing_underrep	14
plot.ROOT	14
reduce_weight	15
ROOT	15
simulated_diabetes_data	17
split_node	18
stratified_kfold	19
summary.characterizing_underrep	19
summary.ROOT	20
train	21
train_single	22

Index**23**

characterize_tree	<i>Fit a shallow decision tree to characterize learned weights w</i>
-------------------	--

Description

Trains a classification tree on the covariates X to predict the binary membership w . This provides an interpretable summary of how the weighted subgroup can be distinguished by X .

Usage

```
characterize_tree(X, w, max_depth = 3)
```

Arguments

X	A data frame of covariates (features).
w	A binary vector (0/1 or a factor with two levels) indicating class membership for each observation (e.g., whether an observation is in the selected subgroup).
max_depth	Integer, the maximum tree depth (default 3).

Details

The tree is grown using the Gini index (classification) and is not pruned (complexity parameter $cp = 0$), relying solely on max_depth to control complexity. This mirrors the default behavior of scikit-learn's `DecisionTreeClassifier(max_depth=...)`. If w is not already a factor, it will be converted internally. The tree's rules can be interpreted to understand which covariates (and what splits) best separate the two classes defined by w .

Value

An `rpart` object representing the fitted decision tree.

characterizing_underrep*Characterize under-represented subgroups (wraps ROOT)*

Description

Combines an RCT (S=1) and a target dataset (S=0), then calls ROOT() to learn a weighted tree that identifies subgroups with different representation in the target population compared to the trial.

Usage

```
characterizing_underrep(
  DataRCT,
  covariateColName_RCT,
  trtColName_RCT,
  outcomeColName_RCT,
  DataTarget,
  covariateColName_TargetData,
  leaf_proba = 0.25,
  seed = 123,
  num_trees = 10,
  vote_threshold = 2/3,
  explore_proba = 0.05,
  feature_est = "Ridge",
  feature_est_args = list(),
  top_k_trees = FALSE,
  k = 10,
  cutoff = "baseline",
  verbose = FALSE,
  global_objective_fn = objective_default,
  keep_threshold = 0.5,
  1X_threshold = NULL
)
```

Arguments

DataRCT	A data frame containing the randomized clinical trial data. Must include treatment, outcome, and covariate columns.
covariateColName_RCT	A character vector of covariate column names in DataRCT.
trtColName_RCT	A character string specifying the treatment column name in DataRCT (0/1).
outcomeColName_RCT	A character string specifying the outcome column name in DataRCT.
DataTarget	A data frame containing the target population data (covariates only).
covariateColName_TargetData	A character vector of covariate column names in DataTarget.
leaf_proba	A numeric value for the "leaf" probability in the ROOT tree growth (default 0.25).
seed	An integer seed for reproducibility (default 123).

num_trees An integer specifying the number of trees to grow (default 10).
 vote_threshold A numeric value (0.5, 1] for the majority vote threshold (default 2/3).
 explore_proba A numeric value for exploration probability (default 0.05).
 feature_est A string ("Ridge", "GBM") or function for feature importance estimation.
 feature_est_args
 A named list of arguments passed to the feature estimator.
 top_k_trees Logical; if TRUE, selects the top k trees instead of using a cutoff.
 k Integer; number of trees to select if top_k_trees is TRUE.
 cutoff A numeric value or "baseline" to determine the Rashomon set cutoff.
 verbose Logical; if TRUE, prints progress and estimand summaries.
 global_objective_fn
 A function(D) -> numeric to minimize (default objective_default).
 keep_threshold Unused; kept for backward compatibility.
 1X_threshold Unused; kept for backward compatibility.

Value

A characterizing_underrep object (S3 list) containing:

root The resulting ROOT object.
 combined The combined data frame (RCT + Target) used for analysis.
 leaf_summary A data frame summarizing the terminal nodes (rules, counts, and labels).

Examples

```

## Not run:
# Load example data
data(diabetes_data)

# Split into Trial (S=1) and Target (S=0)
trial <- subset(diabetes_data, S == 1)
target <- subset(diabetes_data, S == 0)

# Run characterization
res <- characterizing_underrep(
  DataRCT = trial,
  covariateColName_RCT = c("Race_Black", "Sex_Male", "DietYes", "Age45"),
  trtColName_RCT = "Tr",
  outcomeColName_RCT = "Y",
  DataTarget = target,
  covariateColName_TargetData = c("Race_Black", "Sex_Male", "DietYes", "Age45"),
  seed = 123
)
# View Summary
summary(res)

# Plot the annotated tree
plot(res)

## End(Not run)
  
```

choose_feature	<i>Randomly choose a split feature based on provided probabilities</i>
----------------	--

Description

Given a probability distribution over features (and possibly a "leaf" option), selects one feature at random according to those probabilities.

Usage

```
choose_feature(split_feature, depth)
```

Arguments

- split_feature A named numeric vector of feature selection probabilities. Names should correspond to feature IDs (and may include a special "leaf" entry).
- depth Current tree depth (an integer, used for parity with Python implementation but not affecting probabilities in this implementation).

Value

A single feature name (or "leaf") chosen randomly according to the provided probability weights.

Note

The factor $2^{(0*depth/4)}$ present in the code is effectively 1 (no effect on the first element's weight) and is included only for parity with an equivalent Python implementation. All probabilities are normalized to sum to 1 before sampling.

estimate	<i>Compute pseudo-outcome components (a, b) and their product (v)</i>
----------	--

Description

Using the outputs of the nuisance models, computes intermediate values for the treatment effect estimation via inverse probability weighting (IPW) for the Average Treatment Effect (ATE) in trial sample.

Usage

```
estimate(testing_data, outcome, treatment, sample, pi, pi_m, e_m)
```

Arguments

<code>testing_data</code>	A data frame of test data (or evaluation data) containing at least the columns for outcome, treatment, and sample indicators.
<code>outcome</code>	Name of the outcome column in <code>testing_data</code> .
<code>treatment</code>	Name of the treatment column in <code>testing_data</code> (0/1).
<code>sample</code>	Name of the sample indicator column in <code>testing_data</code> (0/1).
<code>pi</code>	Numeric scalar, the estimated $P(S = 1)$ (prevalence) from the training data.
<code>pi_m</code>	A fitted model (e.g., <code>glm</code>) for $P(S = 1 X)$; typically from <code>train()</code> .
<code>e_m</code>	A fitted model (<code>glm</code>) for $P(Tr = 1 X, S = 1)$; typically from <code>train()</code> .

Details

Specifically, it computes:

- a: IPW-adjusted outcome difference, $a_i = S_i \left(\frac{Tr_i Y_i}{p_{t1|x,i}} - \frac{(1-Tr_i)Y_i}{1-p_{t1|x,i}} \right)$.
- b: Overlap weight factor, $b_i = \frac{1}{\ell(X_i)}$, where $\ell(X) = \frac{P(S=1|X)/\pi}{P(S=0|X)/(1-\pi)}$.
- v: The pseudo-outcome, defined as $v_i = a_i \times b_i$.

The predicted probabilities from `pi_m` and `e_m` are constrained to $[1e-8, 1-1e-8]$ to avoid instability (extremely small or large probabilities are clamped). If the provided `pi` is 0 or 1 (indicating no variation in sample inclusion in training), the computation is undefined and an error will be thrown. Ensure that `pi_m` and `e_m` correspond to models trained on compatible data (same covariates) for accurate predictions.

Value

A list with numeric vectors:

<code>v</code>	Pseudo-outcome values for each observation (numeric vector length = <code>nrow(testing_data)</code>).
<code>a</code>	Intermediate "IPW-adjusted outcome" values (same length as <code>v</code>).
<code>b</code>	Overlap weight factors (same length as <code>v</code>).

See Also

`train` for obtaining `pi`, `pi_m`, and `e_m`; `estimate_dml` for cross-fitted estimation.

`estimate_dml`

Cross-fitted estimation of pseudo-outcomes (Double ML)

Description

Trains nuisance models on each training fold and computes pseudo-outcomes on the corresponding test fold, then aggregates results. Returns only what is needed downstream: the pseudo-outcome table and the aligned evaluation data.

Usage

```
estimate_dml(data, outcome, treatment, sample, crossfit = 5)
```

Arguments

<code>data</code>	A data frame containing at least the outcome, treatment, and sample indicator columns.
<code>outcome</code>	Name of the outcome column.
<code>treatment</code>	Name of the treatment column (0/1).
<code>sample</code>	Name of the sample indicator column (0/1).
<code>crossfit</code>	Integer number of folds for cross-fitting (≥ 2).

Value

A list with:

<code>df_v</code>	Data frame with one row per kept observation (indexed by <code>primary_index</code>), containing: <code>te</code> (pseudo-outcome v), <code>a</code> , <code>b</code> , and squared deviations <code>te_sq</code> , <code>a_sq</code> . Only $S==1$ rows with finite values are kept.
<code>data2</code>	Subset of original data corresponding to <code>df_v\$primary_index</code> .

Note

Rows with infinite or undefined weights (e.g., where the predicted propensity scores were 0 or 1) are removed from `df_v` (and the corresponding rows in `data2`). The `primary_index` in `df_v` corresponds to the row index in the original data. Squared deviation columns (`te_sq`, `a_sq`) are centered around the mean of `te` and `a` for the $S==1$ group.

`estimate_dml_single` *Cross-fitted Double ML (single-sample mode)*

Description

Runs K-fold cross-fitting to produce pseudo-outcomes for ATE estimation when no sample-membership indicator is available (or has no variation). On each fold, a treatment propensity model is trained on the training split and used to compute pseudo-outcomes on the test split. Results are combined and centered to create variance proxies.

Usage

```
estimate_dml_single(data, outcome, treatment, crossfit = 5)
```

Arguments

<code>data</code>	A data frame containing <code>outcome</code> , <code>treatment</code> , and covariates.
<code>outcome</code>	Name of the outcome column.
<code>treatment</code>	Name of the binary treatment indicator column (0/1).
<code>crossfit</code>	Integer number of folds for cross-fitting (default 5; must be ≥ 2).

Value

A list with:

- | | |
|--------------------|--|
| <code>df_v</code> | Data frame with one row per kept observation, containing: <code>te</code> (pseudo-outcome v), <code>a</code> , <code>b</code> (all ones), and centered squares <code>te_sq</code> , <code>a_sq</code> , plus <code>primary_index</code> mapping back to the original data rows. |
| <code>data2</code> | Subset of data corresponding to <code>df_v\$primary_index</code> (i.e., rows kept after cross-fitting and finite checks). |

`estimate_single`*Compute single-sample pseudo-outcomes***Description**

Computes the single-sample pseudo-outcome components for ATE-style estimation: $a_i = T_i Y_i / e_i - (1 - T_i) Y_i / (1 - e_i)$, with $v_i = a_i$ and $b_i \equiv 1$. The treatment propensity e_i is predicted from a supplied model.

Usage

```
estimate_single(testing_data, outcome, treatment, e_m)
```

Arguments

- | | |
|---------------------------|---|
| <code>testing_data</code> | A data frame containing at least <code>outcome</code> , <code>treatment</code> , and covariates (the latter are used for prediction). |
| <code>outcome</code> | Name of the outcome column (character). |
| <code>treatment</code> | Name of the binary treatment indicator column (0/1). |
| <code>e_m</code> | A fitted <code>glm(binomial)</code> model for $P(T = 1 X)$; typically the result of train_single . |

Value

A list with numeric vectors of length `nrow(testing_data)`:

- | | |
|----------------|--|
| <code>v</code> | Pseudo-outcome values (equal to <code>a</code> in single-sample mode). |
| <code>a</code> | IPW-adjusted outcome contrast. |
| <code>b</code> | Vector of ones (no sample-overlap weighting in single-sample mode). |

gen_S	<i>Generate sample indicator $S \sim \text{Bernoulli}(\text{plogis}(a))$</i>
-------	---

Description

Generates a binary sample inclusion indicator S for each observation, using a logistic model influenced by a rectangular region in the first two covariates (X_0 and X_1).

Usage

```
gen_S(X, seed = NULL)
```

Arguments

- | | |
|------|--|
| X | A data frame of covariates (must contain at least columns X_0 and X_1). |
| seed | Optional numeric seed for RNG. If provided, <code>set.seed(seed + 1)</code> is invoked for reproducibility. If <code>NULL</code> (default), no specific seed is set. |

Details

The inclusion probability is defined as $p = \text{plogis}(a)$, where $a = 0.25 - 2 * I\{X_0, X_1 \text{ in region } (0.5, 1)\}$. In other words, observations for which both X_0 and X_1 lie in $(0.5, 1)$ have a lower odds of being included (due to a negative contribution in the linear predictor). This mirrors a scenario where a specific region in feature space is under-sampled. If a seed is set, it uses `seed + 1` to differentiate from other generators.

Value

A data frame with a single column S of 0/1 values indicating inclusion (1) or exclusion (0).

gen_T	<i>Generate treatment indicator $\text{Tr} \sim \text{Bernoulli}(pi)$</i>
-------	--

Description

Assigns a treatment indicator for each observation, combining an experimental design for included samples ($S==1$) and an observational assignment for excluded samples ($S==0$).

Usage

```
gen_T(X, S, seed = NULL)
```

Arguments

- | | |
|------|---|
| X | A data frame of covariates. |
| S | A data frame with column S (0/1 indicating sample inclusion for each observation). |
| seed | Optional numeric seed for RNG. If provided, <code>set.seed(seed - 1)</code> is used. Default <code>NULL</code> means no explicit seeding. |

Details

For observations with $S==1$ (in sample), treatment is assigned with probability 0.5 (mimicking a randomized experiment). For those with $S==0$ (out of sample), treatment probability is $\text{plogis}(X0)$, i.e., it increases with the value of covariate $X0$. The overall assignment probability for each observation is $\pi_i = S_i * 0.5 + (1 - S_i) * \text{plogis}(X0_i)$. If a seed is provided, an offset seed - 1 is used to differentiate from other generation steps.

Value

A list with two elements:

Tr	A data frame with a single column Tr (treatment assignments 0/1 for each observation).
pi	A numeric vector of length equal to number of observations, giving the treatment probability used for each observation.

gen_XY

Generate covariates X and potential outcomes (Y0, Y1)

Description

Simulates a regression problem (Friedman #1) and defines a treatment effect. Uses `mlbench.friedman1` to generate X features and a baseline outcome Y0. The treatment potential outcome Y1 is defined as $Y1 = Y0 + \log(Y0 + 1)$, introducing a heterogeneous treatment effect.

Usage

```
gen_XY(n = 1000, seed = NULL)
```

Arguments

n	Integer or numeric. Number of observations to simulate (must be positive).
seed	Optional. Single numeric value for RNG seed. If provided, a global seed is set for reproducibility. If NULL (default), no seed is set (results will vary on each run).

Details

The `mlbench.friedman1` function from the **mlbench** package is used to generate 10 independent continuous features and a baseline outcome Y0 with additive noise. The treatment outcome Y1 is defined by adding a non-linear term $\log(Y0 + 1)$ to the baseline. If a seed is specified, the random number generator state is reset at the start of the function (which affects other random operations).

Value

A list with two components:

X	A data frame of simulated covariates with columns X0, X1, ... up to X(p-1).
Y	A data frame of potential outcomes with columns Y0 (baseline outcome) and Y1 (outcome under treatment).

get_data*Convenience wrapper to generate a full simulated dataset*

Description

Generates covariates, sample inclusion, treatment assignments, and observed outcomes for a specified sample size. This wraps `gen_XY()`, `gen_S()`, and `gen_T()` in sequence.

Usage

```
get_data(n = 1000, seed = NULL)
```

Arguments

<code>n</code>	Integer or numeric. Sample size (number of observations to generate).
<code>seed</code>	Optional base seed for reproducibility. If provided, internal generators use offsets of this seed to ensure independent randomness. Default <code>NULL</code> means no explicit seeding.

Details

This function first generates covariates and potential outcomes with `gen_XY`. It then generates `S` (sample inclusion) and `Tr` (treatment assignment). The observed outcome `Yobs` is computed as $Y_{obs} = Tr * Y1 + (1 - Tr) * Y0$ for each observation.

Value

A list with two components:

<code>data</code>	A data frame of length <code>n</code> containing covariates <code>X0</code> , ..., sample indicator <code>S</code> , treatment indicator <code>Tr</code> , and observed outcome <code>Yobs</code> .
<code>Y</code>	A data frame of length <code>n</code> containing the potential outcomes <code>Y0</code> and <code>Y1</code> for each observation.

Examples

```
sim <- get_data(n = 100, seed = 599)
dim(sim$data)    # should be 100 x (p + 3) columns (p features + S + Tr + Yobs)
head(sim$data$Yobs) # observed outcomes
head(sim$Y)      # potential outcomes corresponding to those observations
```

`loss_from_objective` *Backward/fast-path micro-evaluator adaptor*

Description

Wrap a global objective `global_objective_fn(D)` into a splitter-compatible loss function `loss_fn(val, indices, D)` by evaluating `objective_if` on a temporary copy of `D`.

Usage

```
loss_from_objective(global_objective_fn)
```

Arguments

`global_objective_fn`

Function of one argument `D` returning a numeric scalar to be minimized (e.g., `objective_default`).

Value

A function `loss_fn(val, indices, D)` suitable for use in `ROOT` and `split_node`. It sets `w = val` on `indices` (non-mutating), then returns `global_objective_fn(D)`.

`midpoint` *Compute the midpoint of a numeric vector*

Description

Calculates the midpoint defined as $(\max(X) + \min(X))/2$, ignoring any NA values.

Usage

```
midpoint(X)
```

Arguments

`X` A numeric vector.

Value

A numeric scalar giving the midpoint of the finite values in `X`. If `X` is empty or has no finite values, NA is returned.

<code>objective_default</code>	<i>Default objective: SE proxy of (W)TATE/PATE</i>
--------------------------------	--

Description

Computes $\sqrt{\sum_i vsq_i * w_i / (\sum_i w_i)^2}$. Requires columns vsq and w in D. Minimize this. Supply your own function(D) -> scalar to use a different objective.

Usage

```
objective_default(D)
```

Arguments

D data.frame with at least numeric columns vsq and w.

Value

numeric scalar objective value; Inf if undefined.

<code>objective_if</code>	<i>Helper: evaluate objective after a hypothetical local change</i>
---------------------------	---

Description

Helper: evaluate objective after a hypothetical local change

Usage

```
objective_if(val, indices, D, global_objective_fn)
```

Arguments

val	0/1 assignment to apply
indices	integer or rownames to receive val
D	data.frame used by global_objective_fn
global_objective_fn	function(D)->scalar

Value

numeric scalar objective after the hypothetical change

plot.characterizing_underrep

Plot Under-represented Population Characterization

Description

Visualizes the decision tree derived from the ROOT analysis, highlighting which subgroups are represented ($w=1$) versus underrepresented ($w=0$).

Usage

```
## S3 method for class 'characterizing_underrep'
plot(x, ...)
```

Arguments

- x A characterizing_underrep object.
- ... Additional arguments passed to `rpart.plot::prp()`.

Value

No return value; draws a plot.

plot.ROOT

Plot the ROOT Summary Tree

Description

Visualizes the decision tree that characterizes the weighted subgroup identified by ROOT.

Usage

```
## S3 method for class 'ROOT'
plot(x, ...)
```

Arguments

- x A ROOT object returned by `ROOT()`.
- ... Additional arguments passed to `rpart.plot::prp()`.

Value

No return value; the plot is drawn to the active graphics device.

reduce_weight	<i>Reduce a feature's selection weight by half and renormalize</i>
---------------	--

Description

Lowers the probability weight of a given feature by 50%, and then re-normalizes the entire probability vector.

Usage

```
reduce_weight(fj, split_feature)
```

Arguments

- | | |
|----------------------------|--|
| <code>fj</code> | A feature name (character string) present in the names of <code>split_feature</code> . |
| <code>split_feature</code> | A named numeric vector of probabilities for features (as used in splitting). |

Details

This is typically used when a particular feature split was rejected; the feature's probability is halved to reduce its chance of being chosen again immediately, encouraging exploration of other features. If `fj` is "leaf", its weight is also halved similarly.

Value

A numeric vector of the same length as `split_feature`, giving the updated probabilities that sum to 1.

ROOT	<i>Ensemble of weighted trees (loss/objective-agnostic) and Rashomon selection</i>
------	--

Description

Builds multiple weighted trees, then identifies a "Rashomon set" of top-performing trees and aggregates their weight assignments by majority vote.

Usage

```
ROOT(
  data,
  outcome,
  treatment,
  sample,
  leaf_proba = 0.25,
  seed = NULL,
  num_trees = 10,
  vote_threshold = 2/3,
  explore_proba = 0.05,
  feature_est = "Ridge",
```

```

feature_est_args = list(),
top_k_trees = FALSE,
k = 10,
cutoff = "baseline",
verbose = FALSE,
global_objective_fn = objective_default
)

```

Arguments

<code>data</code>	A data frame containing the dataset. Must include outcome, treatment, and sample indicator columns.
<code>outcome</code>	A character string specifying the name of the outcome column in <code>data</code> .
<code>treatment</code>	A character string specifying the name of the treatment indicator column (0/1) in <code>data</code> .
<code>sample</code>	A character string specifying the name of the sample indicator column (0/1) in <code>data</code> . Use <code>NULL</code> for single-sample SATE mode.
<code>leaf_proba</code>	A numeric value specifying the probability mass for the "leaf" option in each tree (default 0.25).
<code>seed</code>	An integer seed for reproducibility (default <code>NULL</code>).
<code>num_trees</code>	An integer specifying the number of trees to grow in the forest (default 10).
<code>vote_threshold</code>	A numeric value in (0.5, 1] specifying the majority vote threshold for final weight=1 (default 2/3).
<code>explore_proba</code>	A numeric value specifying the probability of exploration at leaves in each tree (default 0.05).
<code>feature_est</code>	A character string ("Ridge", "GBM") or a function(<code>X</code> , <code>y</code> , ...) returning a named, non-negative vector of importances.
<code>feature_est_args</code>	A named list of extra arguments for a user-supplied <code>feature_est</code> function.
<code>top_k_trees</code>	A logical value. If <code>TRUE</code> , select top-k trees by objective; else use <code>cutoff</code> (default <code>FALSE</code>).
<code>k</code>	An integer specifying the number of top trees if <code>top_k_trees = TRUE</code> (default 10).
<code>cutoff</code>	A numeric value or character string "baseline". If <code>top_k_trees = FALSE</code> , this defines the Rashomon set cutoff.
<code>verbose</code>	A logical value. If <code>TRUE</code> , prints 2 lines with (unweighted and weighted) estimate + SE. Default <code>FALSE</code> .
<code>global_objective_fn</code>	A function <code>function(D) -> numeric</code> scoring the entire state (minimize).

Value

An S3 object of class "ROOT" containing:

<code>D_rash</code>	The data frame with weights from the Rashomon set.
<code>f</code>	The summary <code>rpart</code> tree object.
<code>estimate</code>	A list containing the unweighted and weighted estimates.
<code>...</code>	Internal forest structures.

Examples

```
## Not run:  
data(diabetes_data)  
  
# Run ROOT  
res <- ROOT(data = diabetes_data, outcome = "Y", treatment = "Tr", sample = "S")  
  
# Summary of results  
summary(res)  
  
# Plot the characterization tree  
plot(res)  
  
## End(Not run)
```

simulated_diabetes_data

Simulated Diabetes Dataset for Examples

Description

A toy dataset for illustrating ROOT examples and tests.

Usage

```
data(simulated_diabetes_data)
```

Format

A data frame with rows for individuals and the following columns:

Age45 Indicator (0/1): age ≥ 45 .

DietYes Indicator (0/1): on a diet program.

Race_Black Indicator (0/1): race is Black.

S Sample indicator (0/1): 1 = RCT/source, 0 = target.

Sex_Male Indicator (0/1): male.

Tr Treatment assignment (0/1).

Y Observed outcome (numeric or 0/1).

`split_node`*Recursive split builder for weighted tree (internal function)*

Description

Recursively builds a weighted decision tree to optimize a global objective, using an exploration/exploitation trade-off. Internal; used by ROOT().

Usage

```
split_node(
  split_feature,
  X,
  D,
  parent_loss,
  depth,
  explore_proba = 0.05,
  choose_feature_fn = choose_feature,
  reduce_weight_fn = reduce_weight,
  global_objective_fn = objective_default,
  max_depth = 8,
  min_leaf_n = 5,
  log_fn = function(...) {
},
  max_rejects_per_node = 1000
)
```

Arguments

<code>split_feature</code>	Named numeric vector of feature selection probabilities (must include "leaf").
<code>X</code>	Data frame of current observations (includes candidate split feature columns; may include a working copy of weights w).
<code>D</code>	Data frame representing the global state (must include columns w and vsq; row names align to observations).
<code>parent_loss</code>	Numeric, the loss value of the parent node (used to decide if a split improves the objective).
<code>depth</code>	Integer, current tree depth.
<code>explore_proba</code>	Numeric, the probability (between 0 and 1) of flipping the exploit choice at a leaf.
<code>choose_feature_fn</code>	Function to choose next feature (default choose_feature).
<code>reduce_weight_fn</code>	Function to penalize last-tried feature on rejected split (default reduce_weight).
<code>global_objective_fn</code>	Function function(D) -> numeric scoring the entire state.
<code>max_depth</code>	Integer max depth (stop and make leaf at this depth).
<code>min_leaf_n</code>	Integer min rows to attempt a split; else make leaf.
<code>log_fn</code>	Function for logging; default no-op.
<code>max_rejects_per_node</code>	Safety budget of rejected splits before forcing a leaf.

Value

A list representing the (sub)tree; includes updated D and local objective.

stratified_kfold	<i>Stratified K-fold index generator</i>
------------------	--

Description

Splits indices into K folds while preserving the class distribution of a binary factor. This mimics scikit-learn's `StratifiedKFold`, ensuring each fold has a representative ratio of the two classes in S.

Usage

```
stratified_kfold(S, K = 5)
```

Arguments

- | | |
|---|--|
| S | A vector or factor indicating class membership (typically 0/1 or two-class factor) for stratification. |
| K | Integer number of folds ($K \geq 2$). If K is larger than the number of observations, it will be reduced to that number. |

Details

The function deterministically allocates indices to folds by class. For each class in S, indices are cyclically assigned to folds to balance counts. If K == 1, a single fold containing all indices is returned (though typically K should be ≥ 2 for cross-validation).

Value

A list of length K, where each element is an integer vector of row indices assigned to that fold. The union of all folds equals 1:length(S), and folds are roughly equal in size.

summary.characterizing_underrep	<i>Summarize a characterizing_underrep fit</i>
---------------------------------	--

Description

Prints the ROOT summary (un/weighted estimates with standard errors; the *weighted* SE is omitted when a custom `global_objective_fn` was used in `ROOT()`) and a brief overview of terminal rules from the annotated summary tree, if available.

Usage

```
## S3 method for class 'characterizing_underrep'
summary(object, ...)
```

Arguments

- object A characterizing_underrep object.
 ... Unused; included for S3 compatibility.

Details

Delegates core statistics to summary(object\$root); previews up to ten terminal rules when a summary tree exists, and reports plot availability.

Value

object, invisibly.

summary.ROOT	<i>Summarize a ROOT fit</i>
--------------	-----------------------------

Description

Summarizes a ROOT object by reporting the primary estimands and key model diagnostics. The first lines report:

1. the **unweighted** estimate (ATE in RCT for single-sample or TATE for two-sample) and its **standard error (SE)**;
2. the **weighted** estimate (WATE in RCT or WTATE using w_opt) and its **SE** whenever w_opt is effectively *binary* (subset-mean SE); if w_opt is non-binary, the SE is omitted with a note.

Subsequent lines describe the estimand type, number of trees, size of the Rashomon set, presence of a summary tree, covariate count, observation count, baseline loss, selected-tree losses, and the proportion kept by w_opt.

Usage

```
## S3 method for class 'ROOT'
summary(object, ...)
```

Arguments

- object A ROOT object returned by ROOT().
 ... Unused; included for S3 compatibility.

Details

This method prefers pre-computed estimates. If unavailable, it recomputes:

- Unweighted effect as \bar{v} over the analysis set (all rows in single-sample; S = 1 in two-sample);
- Unweighted SE as $\sqrt{\frac{1}{n(n-1)} \sum (v_i - \bar{v})^2} \sqrt{\text{sum}((v - v_{\bar{v}})^2) / (n * (n - 1))}$;
- Weighted effect when w_opt is binary, with $A = \{i : w_i = 1\}$, i.e. $\bar{v}_A = \frac{1}{n_A} \sum_{i \in A} v_i$;
- Weighted SE (WTATE/WATE), for binary w_opt, as $\sqrt{\frac{1}{n_A(n_A-1)} \sum_{i \in A} (v_i - \bar{v}_A)^2} \sqrt{\text{sum}((v[w == 1] - v_{\bar{v}_A})^2) / (n_A * (n_A - 1))}$.

Value

The input object, invisibly. Printed output is a human-readable summary.

train*Train nuisance models for weighting***Description**

Fits models to estimate sampling and treatment propensities on training data. Specifically, it computes:

- `pi`: The prevalence of sample inclusion (estimated as mean of `S`).
- `pi_m`: A logistic regression model for $P(S = 1 | X)$ using all covariates.
- `e_m`: A logistic regression model for $P(Tr = 1 | X, S = 1)$, fit only on the subset where `S==1`.

Usage

```
train(training_data, outcome, treatment, sample)
```

Arguments

<code>training_data</code>	A data frame containing the training dataset.
<code>outcome</code>	Name of the outcome column (typically observed outcome, e.g. "Yobs").
<code>treatment</code>	Name of the treatment indicator column (e.g. "Tr").
<code>sample</code>	Name of the sample inclusion indicator column (e.g. "S").

Details

This function uses simple logistic regression (`glm` with logit link) to estimate the necessary nuisance parameters for weighting. It requires that both `S` and `Tr` have variation (both 0 and 1 must be present in training data); if not, the fitting is not possible and an error is raised. All covariates other than the specified outcome, treatment, and sample columns are used as predictors.

Value

A list with components:

<code>pi</code>	Numeric scalar giving the overall sample inclusion rate $P(S = 1)$ in the training data.
<code>pi_m</code>	A fitted <code>glm</code> model (binomial family) for $P(S = 1 X)$.
<code>e_m</code>	A fitted <code>glm</code> model (binomial family) for $P(Tr = 1 X, S = 1)$.

train_single*Train treatment propensity model (single-sample mode)*

Description

Fits a logistic regression for $P(T = 1 | X)$ on the provided training data. Used by the single-sample Double ML path where no sample-selection model is needed.

Usage

```
train_single(training_data, outcome, treatment)
```

Arguments

- | | |
|----------------------------|--|
| <code>training_data</code> | A data frame containing the outcome, treatment, and covariates. Only <code>treatment</code> and covariates are used for fitting. |
| <code>outcome</code> | Name of the outcome column (character). Present for a consistent signature but not used in this function. |
| <code>treatment</code> | Name of the binary treatment indicator column (0/1). |

Value

A list with one element:

- | | |
|------------------|--|
| <code>e_m</code> | A <code>glm</code> (binomial) object for the treatment propensity model. |
|------------------|--|

Index

* datasets
 simulated_diabetes_data, 17

 characterize_tree, 2
 characterizing_underrep, 3
 choose_feature, 5

 estimate, 5
 estimate_dml, 6, 6
 estimate_dml_single, 7
 estimate_single, 8

 gen_S, 9
 gen_T, 9
 gen_XY, 10
 get_data, 11

 loss_from_objective, 12

 midpoint, 12

 objective_default, 12, 13
 objective_if, 12, 13

 plot.characterizing_underrep, 14
 plot.ROOT, 14

 reduce_weight, 15
 ROOT, 12, 15

 simulated_diabetes_data, 17
 split_node, 12, 18
 stratified_kfold, 19
 summary.characterizing_underrep, 19
 summary.ROOT, 20

 train, 6, 21
 train_single, 8, 22